Cover letter: security and AI evaluation

Zain Dana Harper | Kent, Washington | harperz9.github.io/cv.html | github.com/HarperZ9

Dear hiring team,

I am applying for red-team, security-evaluation and AI-evaluation work. The weak point comes first. I have no professional programming experience and no paid red-team engagement. I spent eleven years in commercial arboriculture and left that trade in June 2026. Everything below was built on my own time and my own systems, with about 13 years of personal programming behind it.

My work is testing whether a check can be fooled. Flywheel, my evaluation harness, is at release 1.9.0. Its behavior check flagged 20 of 20 planted changes and 0 of 50 unchanged pairs. A newer test measures which guards stop destructive tool calls from an agent steered toward a hostile persona. On 28 held-out tasks a schema check caught 1 of 66 harmful calls, and a request-versus-action consistency check caught 64 of 66. That test ran on mocked tools. It is merged, not yet released.

Other maintainers have accepted my fixes. 38 of my engineering pull requests are merged in 31 external repositories, DeepEval and Datasette among them. Another 17 were closed without merge, and the full list is public.

I publish results that failed. My findings page lists 30 results, and 12 of them are nulls. A reviewer who wants to see how I handle being wrong can read those first.

I hold a high school diploma and 42 credentials from 2026, listed with their verification results on my published-work page. I live in Kent, Washington. I will relocate to London, and I can travel to San Francisco regularly.

If the role has a hard requirement, ask me about it directly and I will answer with the evidence or with the gap.

Zain Dana Harper

Written November 13, 2026. Figures and their sources: harperz9.github.io/career/source-ledger.md