



PROJECT TELOS · LET'S HAVE A LOOK

The accountable surface a new era of work **launches from.**

One operation makes any claim, action, or creation *checkable*
— so people and AI can finally build real, high-stakes things
together and verify every step. Accountability isn't the cage;
it's the runway. Security is one domain that takes off from it
— *not the point.*

Zain Dana Harper · github.com/HarperZ9 · it's all public — check it.

THE PROBLEM

AI is capable enough. It just can't be **checked** — so we can't yet build the things that matter with it.

One blind spot strands it: a model *infers* the world but never *senses* it — and sounds equally sure either way. So the wonder stays locked, at two live costs:

AGENTS

act without bounded, witnessed authority. 88% of agent pilots die before production.

OUTPUT

is derivative — recycled, regurgitated. "Slop" was Merriam-Webster's 2025 word; models **provably** memorize their training data.

THE INSIGHT

Accountability is a **systems property**, not a policy. And it's one operation.

PERCEIVE

render any artifact into a witnessed, machine-readable form

CHECK

against a criterion it **did not** author

PROVE

carry a re-checkable proof of the verdict

OR ADMIT

say **UNVERIFIABLE** when you can't

You can't grade your own homework. **That's the anti-hallucination heart.**

THE PRODUCT — SHIPPED, NOT SLIDWARE

A working accountability spine. **Public, tested, run-it-yourself.**

PERCEIVE

coherence-membrane — witnessed perception,
never a screenshot **SHIPPED**

GATE

proof-surface — a default-deny pre-execution gate
· 258 tests **SHIPPED**

THE LOOP

accountable-surface — an MCP server · 201 tests ·
MIT **SHIPPED**

WITNESS

EMET — an external byte-witness, three language
implementations **SHIPPED**

1,100+ tests across the spine. github.com/HarperZ9 — every claim sits beside its test or its honest limit.

THE CATALYST — ONE CORE, EVERY DOMAIN

The incorruptible core is the surface **every** **domain launches from.**

Security checks against *origin*; creativity against *the corpus*; correctness against *a theory* — one loop, different criterion. Everything built on the core inherits the accountability.

DOMAINS

security & provenance · creativity & media ·
research & academia · medicine · simulation & VR
· aerospace, flight, drones · the humanities

Universal as an operation — **demonstrated by instances, extended by binding more.** The falsifiability that keeps it a thesis, not a manifesto.

WEDGES + MARKETS

First into forming categories, on top of validated demand.

CREATIVE — ORIGINALITY

Gen-AI media ~\$2–3B → ~\$20B by 2034.

The "AI slop" pain is *peer-reviewed*;
enterprises already pay for provably-clean
origin (Adobe Firefly).

AGENTS — GOVERNANCE

AI agents ~\$7.8B → ~\$50B by 2030

(~46%); AI red-teaming \$1.3B → \$18.6B
(30%).

The exact categories — originality-verification, accountable-agent-action, provenance — have **no analyst TAM yet**. That's the whitespace. (All figures are analyst estimates.)

WHY NOW

The forcing functions just landed.

REGULATION

EU AI Act — penalties to 7% of global turnover;
high-risk obligations land 2027

STANDARD

MCP became the cross-vendor default (OpenAI,
Google, AWS, Microsoft) and moved under the
Linux Foundation

LAW

the US Copyright Office: purely-AI output isn't
protectable — so a defensible originality claim
unlocks protection

Ethics enforced in the **architecture**, not promised in a deck.

ANTI-CAPTURE

power can't quietly weaponize a system that
witnesses, gates, and journals everything.
Conferral binds the conferrer.

BACKSTOPS

hard limits, enforced: no CSAM · no CBRN/bio ·
own-targets-only · not-for-war

BILATERAL

the standard binds the operator, the model — and
the author. The mandate binds investors too.

**It polices itself from within — binding its
own author. The only thing it costs you
is the power to abuse it. That filters for
mission-aligned capital — a feature.**

THE FOUNDER

Build-and-break range — you make the honest system because you've lived in the dishonest one.

Self-taught, independent since 2023. Compilers, real-time graphics, color science, security, AI accountability — the offensive side and the defensive one, because you learn how a system holds by learning how it breaks.

Execution over credentials: the shipped, tested, adversarially-built work *is* the credential.

H O N E S T S T A G E

Pre-revenue, solo,
early. **The honesty is
the demonstration.**

S T A G E

IP / founder / pre-seed

PRE-REVENUE

V A L U A T I O N

no fantasy figure — quoting one would **disprove
the product**. The number moves with traction, a
second pair of hands, and a credential; not before.

T R A C T I O N

shipped + verified is the traction that exists;
adoption is the next milestone.

THE ASK

Mission-aligned, patient capital — and the introductions a solo founder lacks.

1 · FIRST

non-dilutive — grants & fellowships (no equity,
no mandate-risk)

2 · THEN

open-source adoption — usage becomes the
traction every investor asks for

3 · THEN

a commercial wedge — accountable creativity
and/or agent governance

4 · THEN

mission-aligned equity — from strength, accepting
the mandate

—

Everyone is buying a better tomorrow. This is the surface it gets built on — accountable by construction, so capability can run without being captured. **Security is the floor; the runway is the point.**

`Don't trust it; check it.`

Zain Dana Harper · zaindharper@gmail.com · github.com/HarperZ9 · harperz9.github.io