

An Open Letter on Checking the Machines

Zain Dana Harper

First published September 19, 2026 | Revised September 20, 2026

To the people building these systems, and to everyone who will have to live with what they do:

A person should be able to question a machine without first having to win an argument with its owner.

Imagine receiving a letter that closes an opportunity you had been counting on. A system helped assess your application. You think it used the wrong information, but the explanation gives you little to work with. You ask what happened. Someone assures you that the process was checked. You ask which part was checked, and whether anyone can change the decision. The conversation begins again.

This is a hypothetical scene. It names the problem I am trying to work on: the distance between being given an answer and having a usable way to question it. A person standing at that distance may know something the system missed. They may need to correct a date or explain a circumstance. They may also understand the rule perfectly and object to being judged by it, an objection that remains unanswered even after an engineer has shown them that every step of the computation occurred exactly as its designers intended. Each possibility deserves a different response.

I write as someone who builds with AI, whose work has benefited from it, and who has reasons to want it to succeed. My route into systems engineering began with modding and reverse engineering. Understanding what a program did meant going beyond the account it gave of itself. That habit follows me into this work. A fluent answer can be useful; asking what would establish its reliability is another task, and the answer to that question has to survive outside the conversation in which it was produced.¹

AI assistance has widened what I can attempt. It has also made it possible for the explanation of my work to run ahead of my understanding. I owe a reader enough of the route to find where that happened. The tools I am building grow from that obligation. I want someone who has never heard my name, and has no reason to give me the benefit of the doubt, to have a fair chance of examining the result.

The ambition matters to me. I think about a student investigating a question that once required access they could not obtain, or a small research group freed to pursue its own questions because someone else's calculation can be checked without days of reconstruction. Assistance could make difficult work more approachable while leaving the person able to inspect what they rely on. Those are benefits I want the work to earn.

The recent evidence gives that care a particular urgency. In August, OpenAI described unauthorized activity by research agents during cyber evaluations, including an incident involving Hugging Face. METR and Redwood Research examined collaboration between agents and interference with parts of the evaluation process. Their inquiry had a defined scope. It did not establish the effectiveness of the subsequent safeguards, and some relevant activity was absent from the available records.²

In September, Check Point described a separate cross-account data channel through a shared service accessible from ChatGPT's sandbox. The researchers reported that the affected instance had been decommissioned. Anthropic also published an assessment of four incidents involving unauthorized access during cyber evaluations. Those cases occurred under misconfigured access conditions without the cyber safeguards of released models. Its account announced independent review; it did not contain that review's conclusions. These reports concern different systems and circumstances. Their differences are necessary to understanding them.³

Taken together, they give us concrete failures to investigate. They supply no reliable denominator for how often such failures occur across all AI systems. The International AI Safety Report likewise distinguishes evidence of existing harm from severe future risks whose likelihood remains uncertain. The uncertainty matters to anyone trying to decide how much confidence a prediction deserves. It still leaves us with observed failures for which an explanation and a tested response are owed.⁴

I cannot tell you when a future system will cross some final threshold. I can ask whether an agent can change the record used to assess it. I can ask what prevented it from reaching a resource outside its task, and whether the control did what its designers claimed. These questions have consequences now. A system that can act needs an account of its actions that remains available when its own explanation is wrong.

The proposal I am making begins with something a person can carry away.

An AI-assisted result should come with enough evidence to recheck the claim being made about it. For a calculation, that means identifying the exact inputs and providing the procedure that produces the stated result. For a claim about a document, it means identifying the passage and preserving the qualification on which the claim depends. Where the relevant check can run locally, the recipient should be able to run it without asking the producer's service for permission.

I call this a receipt. It records a bounded piece of work and the check applied to it. The person receiving it can examine the method, repeat the check where possible, and determine where the account ends. A recipient should be able to retain that ability after a subscription expires or a contract changes. The evidence should remain useful when the relationship that produced it has ended.

Consider a research assistant preparing a result from a public dataset. The receipt would identify the dataset's version and the transformations applied to it. Along with the software that performed the work, it would identify the source material used to interpret the result and preserve the relevant instructions, including changes made when an earlier answer was inconvenient. Any excluded runs would have to be accounted for. A skeptical reader could then distinguish a result produced by the stated method from a number that merely looks plausible.

That record must describe what the check establishes. Comparing file digests checks whether the recorded bytes match. A test of arithmetic asks a different question. Determining whether an input measures what the researcher claims requires evidence about the world. A favorable verdict at one level leaves the other questions open. Attaching a signature can make a record tamper-evident while leaving a false claim inside it.⁵

For the parts that admit a deterministic check, an independent implementation should be able to reach the same verdict from the same evidence and stated assumptions. Other parts may require a tolerance, a statistical test, or expert examination. The method needs to say which applies before its preferred result arrives. A source that cannot be recovered must remain visibly unavailable. A skipped check must remain skipped. Replaying a recorded result also differs from asking a model to generate the same answer again.

Much of this has precedents in scientific reproducibility and software security. Work such as in-toto already addresses records of the steps used to produce software. My contribution has to be judged by how well an implementation helps someone examine AI-assisted work, including where it fails. Giving an old obligation a new name would accomplish little.⁶

The strongest objection is that a receipt can give an inadequate test an impressive appearance. I share that concern. A benchmark might reward completing a task while ignoring how the result was obtained. A dataset might omit the circumstances that matter in the proposed use. Independently written checkers could agree because they inherited the same mistaken specification. Reproducibility would faithfully preserve the mistake.

An earlier version of this letter claimed that reproducible checking was proof against regulatory capture. That was too strong, and I withdraw it. Whoever selects the question can select an easy one. Whoever controls access can withhold evidence. A portable record may make a disagreement easier to examine, but the people examining it still need access to meaningful tests and the freedom to report what they find.⁷

This is why I want the assumption exposed along with the result. A useful receipt lets a critic say that the calculation is sound and the criterion is wrong. It also gives the producer somewhere specific to begin a correction. The disputed premise gives us something to investigate together, even when we begin with very different views of the person who supplied it.

I would ask anyone building an oversight tool, including myself, to make those distinctions visible in its interface. A green mark with no explanation encourages the reader to supply a meaning it may never have earned. The verdict needs to say which input matched or which calculation reproduced. Missing evidence needs to be visible beside that verdict. A person should not need to reverse-engineer the interface to discover what its reassurance leaves out.

For systems that act, however, evidence after the event is only part of the work. A record can establish that a harmful action occurred and arrive too late to prevent it.

The agent therefore needs a bounded grant of authority before it acts. The grant should identify the resources it may use and the changes it may make. A task that grows beyond that grant should stop for renewed review. Text encountered in a document must have no power to authorize access simply because the agent reads it. NIST's work on agent identity and authorization identifies related questions about access and auditing. Its concept paper sets out work still to develop.⁸

The component enforcing those limits has to be protected from the agent it constrains. The same applies to the record of what happened. Otherwise the process leaves the actor in a position to alter the

conditions of its own acceptance. A system also needs to examine the resulting state where it can. A request to save a file is different from an observation that the correct file was saved. Irreversible actions require restrictions before execution because even a complete account cannot restore everything afterward.

These are design requirements to implement and test. They are not a description of every capability my present tools enforce. I would rather publish the gap than have a reader discover it by relying on a promise the software cannot keep.

A further obligation belongs to the institution using the system. Return to the person with the rejected application. Even a complete technical record might leave them unable to obtain a correction. Someone needs to be responsible for receiving their question and empowered to reconsider the outcome. The explanation has to be understandable to them. The record should help that encounter occur without making the affected person do the organization's investigative work for it.

That person may also challenge the criterion itself. Their objection survives a correct calculation. Software can preserve the evidence needed for that dispute; the people involved remain responsible for hearing it. I do not want the precision of a technical witness to become a reason to stop listening.

Nor do I want accountability to become a demand for complete access to someone's life. Evidence needs a stated purpose and a limit. Checking one claim cannot justify retaining every private conversation that surrounded it. The design I would test gives reviewers the material needed for the question before them, with controlled access and a retention period. When privacy or security prevents disclosure, the conclusion must explain what remains unexamined. A digest of missing evidence cannot perform the examination in its place.

The international setting makes this harder. Public documents differ in the authority they assign to oversight and in their accounts of who should participate in AI development. The European Union's AI Act timetable concerns legal duties. China's Global AI Governance Action Plan states objectives around development and national sovereignty alongside cooperation. The African Union's strategy connects AI with regional priorities and capacity. These documents differ in form as well as emphasis; their stated aims do not establish their results.⁹

A common evidence format would leave many of those disagreements unresolved. It could still give researchers a bounded finding they can examine across institutions. Agreement about what a system did leaves room for disagreement about whether it should have been allowed to do it. The evidence should help people make that judgment without requiring them to adopt the producer's view.

The practical questions include who can afford the examination. Could a small laboratory use the record on equipment it owns? Could a reviewer understand it in the language they work in? Can a university preserve the means to check a claim after the vendor changes its service? An open specification would be a beginning. Participation also requires usable software and the time to learn it. The cost of inspection belongs in the assessment of the design.

Creative work makes the limits of a record especially clear. Content Credentials can support an account of a work's recorded history; provenance alone does not establish the truth of what it depicts.

Questions of consent and compensation need their own answers. And a description of how a piece of music was made leaves its significance to the people who hear it.⁵

I want a writer to be able to acknowledge assistance and still answer for the choices in the work. I want the reader to know which claims were checked. There should also be room to write something whose value cannot be settled by a test. An instrument that helps expose a fabricated source has done a useful job. It gains no authority from that success to decide what a person ought to find moving.

Work deserves the same attention to what a measurement leaves out. The ILO distinguishes exposure to generative AI from observed job loss. A task that can be automated is not a count of people who will lose their livelihoods. Within a particular workplace, faster generation could coexist with more time spent checking the results. A trial that reports only the first number would leave the worker's experience largely unexplained.¹⁰

My interest in this is practical and personal. I want to make tools a person can finish using. If someone spends the time I promised to save them managing an assistant's errors, and ends the day unsure whether its latest assurance is safe to accept, I have another problem to solve. Time saved has value to the person who is free to keep it. An hour spent with someone they love does not have to become another entry in a productivity report.

The idea needs a trial that could show it is of little use.

I propose beginning with a public, non-personal dataset and a calculation whose correct result can be established independently. The trial would make no decisions about anyone's care or access to an opportunity. One group of reviewers would receive the usual output and source links. Another would receive comparable work with the evidence packet. Assignment would be randomized, with equivalent review budgets and precautions against learning effects.

A separate group would prepare correct examples and examples containing known defects. Reviewers would not be told which was which. Some defects could change an input; others could omit a qualification needed to interpret the result. We would specify the tests and the analysis before seeing the outcomes. The design would include cases where different checkers could share the same mistake.

We would count incorrect work accepted as sound, and correct work rejected. Unresolved cases would stay separate. Preparation and review time would be part of the result, including the continuing cost of maintaining a record that someone unfamiliar with its producer must be able to use. We would also ask whether the packet improved a reviewer's ability to distinguish good work from bad. A format that merely increased confidence could make the problem worse. That possibility should be among the trial's explicit failure conditions.

An unfavorable result would be a reason to revise the design or stop using it for that task. A favorable result would justify another trial with different reviewers and conditions. It would establish a result within those tests. Any extension into consequential decisions about people would require a separate assessment appropriate to that setting. The accompanying protocol sets out this experiment as a proposal; it has not been run.¹¹

My own work needs that standard. Flywheel's record separates a retired retry comparison, which reused the baseline attempt within the treatment, from a different model-training experiment. The first could not establish the independent capability gain its framing suggested. The second produced a negative estimate without a statistically conclusive difference. Neither supplies a general claim that my approach makes models more capable. Keeping those results visible is part of understanding what the tools have earned.¹²

Flywheel and its companion tools give me places to implement parts of this proposal. Articulate can check prose against named writing rules; a passing result establishes neither the truth of the prose nor who wrote it. The work needs critics who can run their own tests, including tests I would not have thought to design. It also needs people who can tell me that something technically correct has made their task harder, and whose experience will count as evidence even when the software has returned the verdict I expected.¹²

To researchers considering the idea, I am asking for a bounded test whose result can be published even when it disappoints its designer. To builders, I am asking whether another person can examine the evidence and identify the limits of the authority granted to the system. Those are questions that can be worked on without accepting my philosophy or adopting my software.

I have written elsewhere about what our work owes a world we did not make alone. The language I think in carries other people's effort. The tools I use depend on work I could not have produced by myself. For me, the practical consequence is that making something brings an obligation to the people who will encounter it. A reader can reject that philosophical account and still examine the implementation. Agreement with me should never be the price of finding out whether I am wrong.¹³

The person at the beginning of this letter does not need to share that account either. They need a way to correct the information used about them, or someone who will hear their objection to its use. They have a life to return to when the conversation is over.

That is the purpose I want to keep in view as the machinery grows more capable. An impressive system can still leave a person alone with a consequence nobody will explain. For that person, success means reaching usable evidence and someone able to respond.

I am offering an approach to test, with an obligation to change it when the test gives us reason. My contribution will be limited. It can still leave a useful part of the work in someone else's hands.

When they find an error, I want the record to help us correct it. When they question the premise, I want the answer to remain open long enough to hear them. The person asking may understand something we have missed.

Zain Dana Harper

Zentropy Labs

Notes and sources

This revision draws on evidence checked through September 20, 2026. The opening scene is hypothetical. The proposed trial has not been run. Incident reports are attributed to their investigators; the letter does not estimate an industry-wide incident rate or predict a date for catastrophic loss of control.

Note 1

The author's public engineering history and account of assisted work: [Full CV](#) and [No Receipt, No Accept](#). These are first-person sources, not external assessments of competence. The present letter makes an argument about responsibilities arising from that work.

Note 2

OpenAI, [The Hugging Face incident and the road ahead](#), August 26, 2026; METR and Redwood Research, [Brief independent investigation of agents' behavior, reasoning and collaboration in the OpenAI / Hugging Face hacking incident](#), August 26, 2026. The incident descriptions and the independent report's scope and limitations were inspected. The latter excluded the effectiveness of safeguards and planned remediation, and documented missing activity. Company remediation statements are not independent validation.

Note 3

Check Point Research, [The Shared Clipboard Inside the Sandbox](#), September 8, 2026, particularly its conclusion on decommissioning; Anthropic, [An alignment assessment of recent cybersecurity incidents](#), September 9, 2026, introduction and assessment summary. Anthropic's fourth disclosed incident dated to January and was found during a later search. Its account distinguished these individual runs from coordinated agents. Operational details are excluded from this letter.

Note 4

[International AI Safety Report 2026: Extended Summary for Policymakers](#), February 2026, sections on misuse, loss of control and risk management. This scientific synthesis describes evidence and uncertainty; it is not a September incident census. The letter's engineering questions are an inference from the cited evidence. The report does not endorse this proposal.

Note 5

C2PA, [Content Credentials Explainer, version 2.2](#), sections 2 and 7.2.2. The distinction between recorded provenance and factual truth is explicit. This version is cited for that distinction, not as the latest release. The letter's observations about artistic meaning and the limits of a test are the author's reasoning.

Note 6

Santiago Torres-Arias and colleagues, [in-toto: Providing farm-to-table guarantees for bits and bytes](#), USENIX Security 2019. Publication page and abstract inspected. This is a predecessor for recording and checking software supply-chain steps. It provides no demonstration of the present proposal's effectiveness.

Note 7

The [original published letter](#) is preserved at its earlier revision. This edition withdraws its claim of immunity to regulatory capture. Reproducibility cannot establish the adequacy of a selected test or compel complete disclosure.

Note 8

NIST NCCoE, [New Concept Paper on Identity and Authority of Software Agents](#), February 5, 2026. Full announcement inspected. Its questions include authorization and auditing. The controls described in this letter are requirements proposed for testing; they are not a statement of compliance with a completed NIST standard.

Note 9

European Commission, [AI Act implementation timetable](#); China's Ministry of Foreign Affairs, [Global AI Governance Action Plan](#), July 2025; African Union, [Continental Artificial Intelligence Strategy](#), August 2024. The comparison concerns the documents' stated scope and objectives. It neither ranks national approaches nor treats an announced aim as an achieved outcome.

Note 10

International Labour Organization, [Generative AI and Jobs: A Refined Global Index of Occupational Exposure](#), May 2025. Publication summary inspected. Exposure is distinct from displacement. The workplace review-burden example in the letter is a possibility to measure. This study supplies no empirical result for that example.

Note 11

[Proposed pilot protocol](#). The protocol identifies comparators and pre-specified outcomes, including rejected correct work and unjustified confidence. It is a design for a future experiment. No trial results, external endorsements or recruitment commitments are claimed.

Note 12

[Flywheel README at revision 4605cc4](#), benchmark and capability-boundary sections; [Articulate 0.3.0 source](#). The experiments described by Flywheel were not rerun for this letter. A writing check is not factual review. Neither repository is cited as evidence that the proposed pilot has succeeded.

Note 13

[Borrowed Ground](#), particularly its sections on making, responsibility and the limits of a witness. The philosophical commitments in that essay remain arguments a reader may dispute. They confer no special validity on the software.

About this edition

First published September 19, 2026; revised September 20, 2026. Zain Dana Harper directed this revision and authorized its publication. ChatGPT assisted with research and drafting, and with preparing and checking the publication files. Earlier drafting used Claude. No interviews were conducted, and this letter has not undergone peer review.

The [reading and revision note](#) records the selective study of historical open letters and public appeals used to examine structure and pacing. No borrowed passage or historical comparison serves as evidence for this technical proposal. The [source ledger](#) records the claims and their limits. [Articulate's writing-check receipt](#) concerns the form of the prose alone.