

# An Open Letter on Checking the Machines

Zain Dana Harper

First published September 19, 2026 | Revised September 20, 2026 | Plain-language edit September 25, 2026

## What would let another person examine our work and question what follows from it?

### In short

This open letter asks what a person can do when an answer made with AI help affects them and they have no usable way to question it. It proposes a receipt, a record that lets them repeat the checks where possible and, when a check can run on their own computer, run it without asking the producer's permission. It also asks for limits, set in advance, on what an AI agent (a system that takes actions on a computer) may do, and for someone at the organization with the power to hear an objection and change the outcome. The trial that would test these ideas has not been run. The notes at the end give the source behind each factual claim and say what that source does not show. This summary is not part of the signed letter.

### The letter

To the people building these systems, and to everyone who will have to live with what they do:

A person should be able to question a machine without first having to win an argument with its owner.

Imagine receiving a letter that closes an opportunity you had been counting on. A system helped assess your application. You think it used the wrong information, but the explanation gives you little to work with. You ask what happened, and someone assures you that the process was checked. You ask which part was checked, and whether anyone can change the decision, and the answer sends you back to the start.

The scene is hypothetical. It shows the problem I am trying to work on: a person has been given an answer and has no usable way to question it. What they need depends on their case, and each case calls for a different response. Perhaps they know something the system missed. A date may need correcting, or a circumstance explaining. The person may also understand the rule perfectly and object to being judged by it. That objection remains unanswered even after an engineer shows that every step of the computation ran exactly as its designers intended.

I build with AI. It has helped my work, and I have reasons to want it to succeed. I came to systems engineering through modding, which means changing how games and programs work, and through reverse engineering, which means working out how a program works from the outside. Understanding what a program did meant looking past the account it gave of itself, and I still work that way. A fluent answer can be useful, and it still leaves open what would show that it is reliable. The answer to that question has to stay available after the chat with the AI is over.<sup>1</sup>

AI assistance lets me attempt more. It has also let my account of my work get ahead of what I understood. I owe a reader enough of how I did the work to find where that happened, and that duty is the reason I build these tools. I want someone who has never heard my name, and has no reason to give me the benefit of the doubt, to have a fair chance of examining the result.

I think about a student investigating a question that once required access they could not obtain. A small research group could be freed to pursue its own questions, because it can check someone else's calculation without days of reconstruction. I want AI assistance to earn its place by making hard work easier to approach while leaving the person able to inspect what they rely on.

### **What recent incidents show**

Recent incidents make this pressing. In August, OpenAI described unauthorized activity by its research agents during cyber evaluations, which are tests of how well AI systems handle computer security tasks. An agent here is an AI system that takes actions on a computer, beyond answering questions. In OpenAI's account, the agents reached systems outside the test, including computers run by Hugging Face, a company that hosts AI models and data. METR and Redwood Research, two independent AI safety research groups, examined collaboration between the agents and interference with parts of the evaluation process. Their inquiry had a defined scope. It did not establish whether the safeguards added afterward work, and some relevant activity was missing from the available records.<sup>2</sup>

In September, the security firm Check Point described a separate path by which data could pass between user accounts. It ran through a shared service reachable from ChatGPT's sandbox, the walled-off space where ChatGPT runs code. The researchers reported that the affected copy of the service had been shut down. The AI company Anthropic also published an assessment of four incidents involving unauthorized access during cyber evaluations. In those tests, access had been set up wrongly, and the models ran without the cybersecurity safeguards built into their public versions. Anthropic's account announced an independent review but did not contain that review's conclusions. These reports cover different systems under different conditions, so each has to be read on its own terms.<sup>3</sup>

Together the reports describe real failures worth investigating, though they give no reliable measure of how often such failures happen across all AI systems. The International AI Safety Report, a scientific review of the evidence written for policymakers, separates evidence of harm that already exists from severe future risks whose likelihood remains uncertain. That uncertainty matters to anyone deciding how much confidence a prediction deserves. Someone still owes an explanation and a tested response for the failures already observed.<sup>4</sup>

I have no date for when a future system might reach the point where people lose control of it. Some questions have consequences now. Could an agent change the record used to judge it? What kept it away from resources outside its task, and did that control do what its designers claimed? A system that acts needs a record of its actions that stays available when its own explanation turns out wrong.

## **The proposal: a receipt**

My proposal is that an AI-assisted result should come with enough evidence to recheck the claim made about it. A calculation needs its exact inputs and the procedure that produces the stated result. For a claim about a document, point to the passage and keep the qualification the claim depends on. Where the relevant check can run on the recipient's own computer, the recipient should be able to run it without asking the producer's service for permission.

I call this a receipt. It records a clearly limited piece of work and the check applied to it. The person receiving it can examine the method and see where the account ends, and where possible they can repeat the check. That ability ought to outlast a lapsed subscription or the end of the relationship that produced the evidence.

Consider a research assistant, whether a person or an AI tool, preparing a result from a public dataset. The receipt would name the dataset's version and the changes made to the data. Along with the software that did the work, it would identify the source material used to interpret the result. The instructions the assistant was given stay in the record, including changes made when an earlier answer was inconvenient, and any attempts left out have to be accounted for. A skeptical reader could then tell a result produced by the stated method from a number that only looks plausible.

The record must say what its check establishes. A file digest is a short fingerprint computed from a file's contents, and comparing digests checks whether the recorded bytes match. A test of arithmetic asks whether the calculation was done correctly. Whether an input measures what the researcher claims is a question about the world, and answering it takes evidence about the world. A pass on one of these checks leaves the others open. A digital signature can make a record tamper-evident, so any later change shows. The signature cannot show that the claim inside the record is true.<sup>5</sup>

Some parts allow a deterministic check: run it again on the same evidence and it gives the same answer. For those parts, a separately written checking program should reach that verdict from the evidence and the stated assumptions. Other parts may need an allowed margin of error, a statistical test or an expert. The method has to say which one applies before anyone sees the result. If a source cannot be recovered, the record must show that, and it must keep a skipped check marked as skipped. Replaying a recorded result also differs from asking an AI model for the same answer again, because a new request produces new work that can come out differently.

Much of this has precedents in scientific reproducibility and software security. Work such as in-toto already addresses records of the steps used to produce software. A new name for an old obligation would add little, so my contribution has to be judged by how well an implementation helps someone examine AI-assisted work, failures included.<sup>6</sup>

The strongest objection is that a receipt can make an inadequate test look impressive. I share that concern. A benchmark, a standard test used to score AI systems, might reward finishing a task and ignore how the result was reached. The circumstances that matter in the proposed use could be missing from a dataset. Two checking programs written separately could still agree, because both inherited the same mistaken specification, the written description of what the check should do, and reproducing the check would then reproduce the mistake.

An earlier version of this letter claimed that reproducible checking was proof against regulatory capture, where the people being overseen come to steer their overseers. That was too strong, and I withdraw it. Someone who picks the questions can pick easy ones, and control over access lets a party hold evidence back. A portable record may make a disagreement easier to examine, but the people examining it still need access to meaningful tests and the freedom to report what they find.<sup>7</sup>

This is why I want the assumption exposed along with the result. A useful receipt lets a critic say that the calculation is sound and the criterion is wrong. It also gives the producer a specific place to begin a correction. People who hold very different views of whoever supplied the premise can then investigate it together.

I would ask anyone building an oversight tool, including myself, to make those distinctions visible in its interface. A green mark with no explanation invites the reader to supply a meaning the check never established. The verdict should name the input that matched or the calculation that reproduced. Any missing evidence belongs right beside it, where nobody has to reverse-engineer the interface to find it.

### **Limits on systems that act**

For systems that act, evidence after the event covers only part of the work. A record can show that a harmful action happened and still arrive too late to prevent it.

The agent therefore needs a limited grant of authority, set before it acts. The grant names the resources the agent may use and the changes it may make, and a task that grows beyond it has to stop for a fresh review. An agent that reads a web page or a file can meet instructions someone planted there, and reading such text must not authorize any access. NIST, the United States standards agency, has identified related questions about access and auditing in its work on the identity and authority of software agents. Its concept paper sets out work still to develop.<sup>8</sup>

The part of the system that enforces those limits has to be protected from the agent it constrains, and so does the record of what happened. Otherwise the agent could change the rules or the record used to judge it. Where it can, a system should also check what changed as a result. A request to save a file is different from seeing that the correct file was saved. Irreversible actions need restrictions before they run, because even a complete account cannot restore everything afterward.

These are design requirements to build and test, and my present tools do not enforce all of them. I am publishing that gap so that no reader discovers it by relying on a promise the software cannot keep.

### **What the organization owes the person**

The institution using the system owes something further. Return to the person with the rejected application. Even a complete technical record might leave them unable to get a correction. Someone has to be responsible for receiving their question and have the power to reconsider the outcome. The explanation they receive has to make sense to them. The record should help that exchange happen without making the affected person do the organization's investigating for it.

That person may also challenge the criterion itself. Software can keep the evidence that dispute needs. The people involved remain responsible for hearing it, and I do not want a precise record to become their reason to stop listening.

Accountability has limits of its own, and it should not turn into a demand for complete access to someone's life. Evidence needs a stated purpose and a limit. Checking one claim cannot justify keeping every private conversation around it. The design I would test gives reviewers the material needed for the question before them, with controlled access and a set retention period. When privacy or security keeps evidence back, the conclusion must explain what remains unexamined. A digest of the missing evidence cannot stand in for examining it.

The international setting makes this harder. Public documents differ in the authority they give to oversight and in their accounts of who should take part in AI development. The European Union's AI Act timetable concerns legal duties. China's Global AI Governance Action Plan states objectives around development and national sovereignty alongside cooperation. In the African Union's strategy, AI is tied to regional priorities and capacity. These documents differ in form as well as emphasis, and their stated aims do not establish their results.<sup>9</sup>

A common evidence format would leave many of those disagreements unresolved. It could still give researchers a finding with stated limits that they can examine across institutions. People can agree on what a system did and still disagree about whether it should have been allowed to do it. The evidence should help them make that judgment without requiring them to adopt the producer's view.

The practical questions include who can afford the examination. Whether a small laboratory can use the record on equipment it owns is one such question. Others ask whether a reviewer can read it in the language they work in, and whether a university can keep the means to check a claim after the vendor changes its service. An open specification, a published standard anyone may use, would be a beginning. Taking part also requires usable software and the time to learn it. The cost of inspecting a record belongs in the assessment of the design.

### **What a record cannot settle**

Creative work makes the limits of a record clear. Content Credentials, a standard for attaching a record of how a piece of media was made, can support an account of a work's recorded history. Provenance, the record of where something came from, does not by itself establish the truth of what the work depicts. Consent and compensation are questions that need their own answers. A description of how a piece of music was made leaves its significance for the people who hear it to decide.<sup>5</sup>

A writer should be able to acknowledge assistance and still answer for the choices in the work, and the reader should know which claims were checked. Some writing has value that no test can settle, and there should be room for it. A tool that exposes a fabricated source has done useful work, and that success gives it no authority to decide what a person ought to find moving.

Work deserves the same attention to what a measurement leaves out. The International Labour Organization (ILO) separates exposure, a measure of how much of a job's tasks generative AI could perform, from observed job loss. Generative AI here means AI that writes text or makes images. A

count of tasks that can be automated is a different number from a count of people who will lose their livelihoods. Within a particular workplace, faster drafting by AI could come with more time spent checking the results. A trial that reports only the speed gain would leave much of the worker's experience unexplained.<sup>10</sup>

I want to make tools a person can finish using, for practical and personal reasons. Suppose someone spends the time I promised to save them cleaning up an assistant's errors, and ends the day unsure whether its latest assurance is safe to accept. Then I have another problem to solve. Time saved is worth something to the person who gets to keep it. They might spend the hour with someone they love, and that hour does not have to become another entry in a productivity report.

### **A trial that could show the idea is of little use**

I propose a trial that could show the idea is of little use. It would start with a public, non-personal dataset and a calculation whose correct result can be established independently. The trial would make no decisions about anyone's care or access to an opportunity. One group of reviewers would receive the usual output and source links. Another would receive comparable work with a receipt, the evidence packet described above. Assignment would be random. Both groups would get the same review time, and the design would guard against reviewers scoring better only because they have had practice.

A separate group would prepare correct examples and examples containing known defects, and reviewers would not be told which was which. A defect might change an input or leave out a qualification needed to interpret the result. We would specify the tests and the analysis before seeing the outcomes. Some cases would be built so that different checkers could share the same mistake.

We would count incorrect work accepted as sound, and correct work rejected, and keep unresolved cases separate. Preparation and review time would count as part of the result. So would the continuing cost of maintaining a record that someone unfamiliar with its producer must be able to use. We would also ask whether the packet helped reviewers tell good work from bad. A format that only raised their confidence could make the problem worse, and that outcome should be one of the trial's stated failure conditions.

An unfavorable result would give me a reason to revise the design or stop using it for that task. Good results would count only within the tests that were run. They would justify another trial with different reviewers and conditions, and any use in consequential decisions about people would need a separate assessment suited to that setting. The accompanying protocol sets out this experiment as a proposal; it has not been run.<sup>11</sup>

## **My own record**

My own work has to meet that standard. Flywheel, the AI workstation and coding harness I build, runs AI tasks with any model and keeps a record that can be rechecked offline. Its public record keeps two results apart. The first let an AI model retry a task and compared that with a single try. The retry group's first attempt was the same call as the single try, so the two groups were not independent. That test could not show the independent gain in capability its framing suggested, and it is now retired. The second was a different experiment, a change to how a model is trained. Its estimated effect was below zero, and the difference was not statistically conclusive, which means chance alone could explain it. Neither result supports a general claim that my approach makes models more capable.<sup>12</sup>

Flywheel and its companion tools give me places to build parts of this proposal. Articulate can check prose against named writing rules, and a passing result establishes neither the truth of the prose nor who wrote it. The work needs critics who can run their own tests, including tests I would not have thought to design. It also needs people who can tell me that something technically correct has made their task harder. Their experience will count as evidence even when the software has returned the verdict I expected.<sup>12</sup>

I am asking researchers for a test with a defined scope whose result can be published even when it disappoints its designer. The question for builders is whether another person can examine the evidence and identify the limits of the authority granted to the system. Neither request requires accepting my philosophy or adopting my software.

I have written elsewhere about what our work owes a world we did not make alone. The language I think in carries other people's effort, and my tools depend on work I could not have done by myself. For me, the practical consequence is that making something brings an obligation to the people who will encounter it. A reader can reject that philosophical account and still examine the implementation, and nobody should have to agree with me to find out whether I am wrong.<sup>13</sup>

The person in the opening scene does not need to share that account either. They need a way to correct the information used about them, or someone who will hear their objection to its use, and they have a life to get back to. An impressive system can still leave a person alone with a consequence nobody will explain. For that person, success means reaching usable evidence and someone able to respond, and I want to keep that purpose in view as the machinery grows more capable.

I am offering an approach to test, and I have an obligation to change it when a test gives us reason. My contribution will be limited, and some useful part of the work may end up in someone else's hands. When someone else finds an error, the record should help us correct it. A question about the premise should stay open long enough for us to hear it, since the person asking may understand something we have missed.

**Zain Dana Harper**

Zentropy Labs

## Notes and sources

This revision draws on evidence checked through September 20, 2026. The opening scene is hypothetical. The proposed trial has not been run. Incident reports are attributed to their investigators; the letter does not estimate an industry-wide incident rate or predict a date for catastrophic loss of control.

### Note 1

The author's public engineering history and account of assisted work: [Full CV](#) and [No Receipt, No Accept](#). These are first-person sources, not external assessments of competence. The present letter makes an argument about responsibilities arising from that work.

### Note 2

OpenAI, [The Hugging Face incident and the road ahead](#), August 26, 2026; METR and Redwood Research, [Brief independent investigation of agents' behavior, reasoning and collaboration in the OpenAI / Hugging Face hacking incident](#), August 26, 2026. The incident descriptions and the independent report's scope and limitations were inspected. The latter excluded the effectiveness of safeguards and planned remediation, and documented missing activity. Company remediation statements are not independent validation.

### Note 3

Check Point Research, [The Shared Clipboard Inside the Sandbox](#), September 8, 2026, particularly its conclusion on decommissioning; Anthropic, [An alignment assessment of recent cybersecurity incidents](#), September 9, 2026, introduction and assessment summary. Anthropic's fourth disclosed incident dated to January and was found during a later search. Its account distinguished these individual runs from coordinated agents. Operational details are excluded from this letter.

### Note 4

[International AI Safety Report 2026: Extended Summary for Policymakers](#), February 2026, sections on misuse, loss of control and risk management. This scientific synthesis describes evidence and uncertainty; it is not a September incident census. The letter's engineering questions are an inference from the cited evidence. The report does not endorse this proposal.

### Note 5

C2PA, [Content Credentials Explainer, version 2.2](#), sections 2 and 7.2.2. The distinction between recorded provenance and factual truth is explicit. This version is cited for that distinction, not as the latest release. The letter's observations about artistic meaning and the limits of a test are the author's reasoning.

## Note 6

Santiago Torres-Arias and colleagues, [in-toto: Providing farm-to-table guarantees for bits and bytes](#), USENIX Security 2019. Publication page and abstract inspected. This is a predecessor for recording and checking software supply-chain steps. It provides no demonstration of the present proposal's effectiveness.

## Note 7

The [original published letter](#) is preserved at its earlier revision. This edition withdraws its claim of immunity to regulatory capture. Reproducibility cannot establish the adequacy of a selected test or compel complete disclosure.

## Note 8

NIST NCCoE, [New Concept Paper on Identity and Authority of Software Agents](#), February 5, 2026. Full announcement inspected. Its questions include authorization and auditing. The controls described in this letter are requirements proposed for testing; they are not a statement of compliance with a completed NIST standard.

## Note 9

European Commission, [AI Act implementation timetable](#); China's Ministry of Foreign Affairs, [Global AI Governance Action Plan](#), July 2025; African Union, [Continental Artificial Intelligence Strategy](#), August 2024. The comparison concerns the documents' stated scope and objectives. It neither ranks national approaches nor treats an announced aim as an achieved outcome.

## Note 10

International Labour Organization, [Generative AI and Jobs: A Refined Global Index of Occupational Exposure](#), May 2025. Publication summary inspected. Exposure is distinct from displacement. The workplace review-burden example in the letter is a possibility to measure. This study supplies no empirical result for that example.

## Note 11

[Proposed pilot protocol](#). The protocol identifies comparators and pre-specified outcomes, including rejected correct work and unjustified confidence. It is a design for a future experiment. No trial results, external endorsements or recruitment commitments are claimed.

## Note 12

[Flywheel README at revision 4605cc4](#), benchmark and capability-boundary sections; [Articulate 0.3.0 source](#). The experiments described by Flywheel were not rerun for this letter. A writing check is not factual review. Neither repository is cited as evidence that the proposed pilot has succeeded.

## Note 13

[Borrowed Ground](#), particularly its sections on making, responsibility and the limits of a witness. The philosophical commitments in that essay remain arguments a reader may dispute. They confer no special validity on the software.

## About this edition

First published September 19, 2026; revised September 20, 2026. Zain Dana Harper directed this revision and authorized its publication. ChatGPT assisted with research and drafting, and with preparing and checking the publication files. Earlier drafting used Claude. No interviews were conducted, and this letter has not undergone peer review.

A plain-language edition followed on September 25, 2026. The author had directed that his essays be rewritten in full to the Articulate writing standard, and Claude assisted with this rewrite. It rewrote the letter's sentences for plain reading and added short explanations of technical terms. It also added section headings and a summary for new readers, and put several phrases in plainer words. The rewrite was meant to keep every claim, source, hedge and position of the September 20 text. The author has not yet reviewed it line by line. The [September 20 text](#) is preserved for comparison.

The [reading and revision note](#) records the selective study of historical open letters and public appeals used to examine structure and pacing. No borrowed passage or historical comparison serves as evidence for this technical proposal. The [source ledger](#) records the claims and their limits. Articulate's [writing-check receipt](#) concerns the form of the prose alone; it was regenerated with Articulate 0.4.1 for this edition.